

DATA 621 Assignment 2

Joby John, Zachary Herold, Jun Pan

March 11, 2019

Set working environment

Load data

```
output <- read.csv("https://raw.githubusercontent.com/johnpannyc/group-1-data-621-assignment-2/master/classification-output-data.csv")
```

1. DATA EXPLORATION

```
glimpse(output)
```

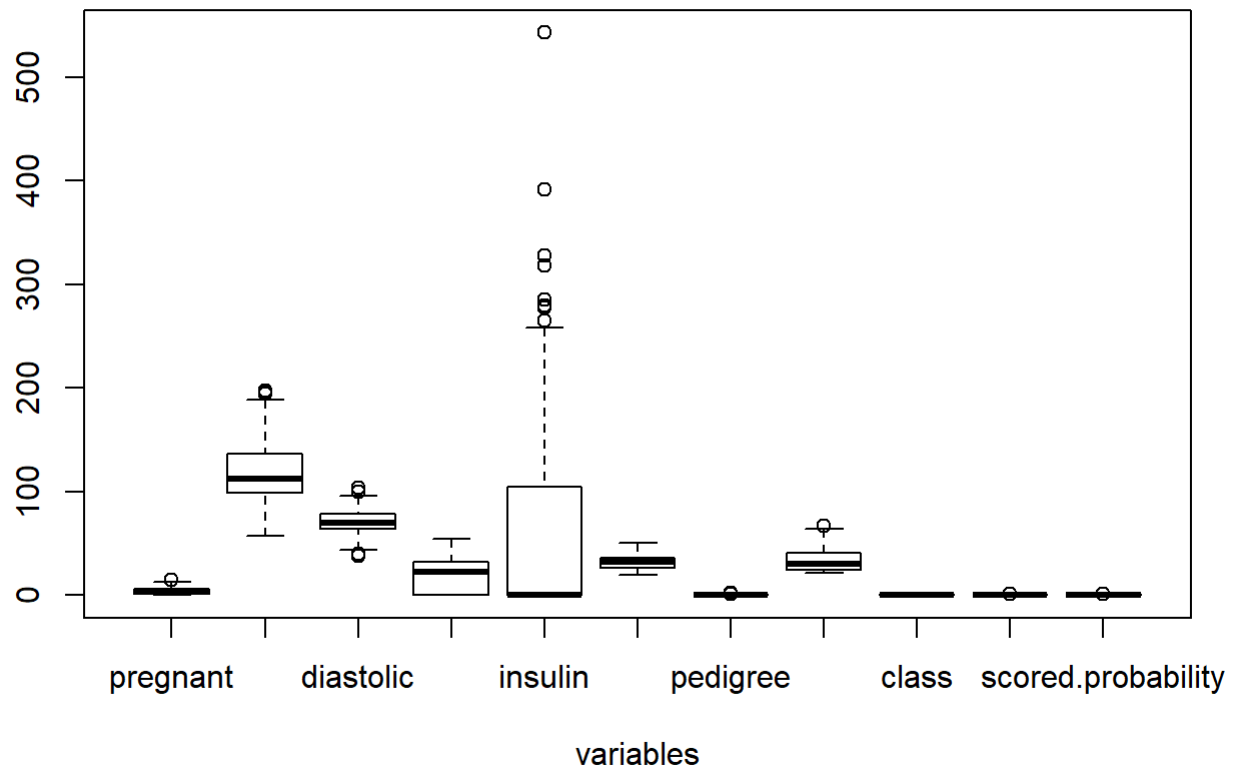
```
## Observations: 181
## Variables: 11
## $ pregnant      <int> 7, 2, 3, 1, 4, 1, 9, 8, 1, 2, 5, 5, 13, 0, ...
## $ glucose       <int> 124, 122, 107, 91, 83, 100, 89, 120, 79, 12...
## $ diastolic     <int> 70, 76, 62, 64, 86, 74, 62, 78, 60, 48, 78,...
## $ skinfold      <int> 33, 27, 13, 24, 19, 12, 0, 0, 42, 32, 30, 4...
## $ insulin       <int> 215, 200, 48, 0, 0, 46, 0, 0, 48, 165, 0, 7...
## $ bmi           <dbl> 25.5, 35.9, 22.9, 29.2, 29.3, 19.5, 22.5, 2...
## $ pedigree      <dbl> 0.161, 0.483, 0.678, 0.192, 0.317, 0.149, 0...
## $ age           <int> 37, 26, 23, 21, 34, 28, 33, 64, 23, 26, 37,...
## $ class         <int> 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1...
## $ scored.class  <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1...
## $ scored.probability <dbl> 0.32845226, 0.27319044, 0.10966039, 0.05599...
```

classification-output-data.csv file contains 181 observations of 11 variables. Three variables will be considered for this report - class (actual class for the observation), scored.class (predicted class for the observation), and scored.probability (predicted probability of success for the observation).

```
summary(output)
```


```
##      pregnant      glucose      diastolic      skinfold
## Min.   : 0.000   Min.   : 57.0   Min.   : 38.0   Min.   : 0.0
## 1st Qu.: 1.000   1st Qu.: 99.0   1st Qu.: 64.0   1st Qu.: 0.0
## Median : 3.000   Median :112.0   Median : 70.0   Median :22.0
## Mean   : 3.862   Mean   :118.3   Mean   : 71.7   Mean   :19.8
## 3rd Qu.: 6.000   3rd Qu.:136.0   3rd Qu.: 78.0   3rd Qu.:32.0
## Max.   :15.000   Max.   :197.0   Max.   :104.0   Max.   :54.0
##      insulin      bmi      pedigree      age
## Min.   : 0.00   Min.   :19.40   Min.   :0.0850   Min.   :21.00
## 1st Qu.: 0.00   1st Qu.:26.30   1st Qu.:0.2570   1st Qu.:24.00
## Median : 0.00   Median :31.60   Median :0.3910   Median :30.00
## Mean   : 63.77   Mean   :31.58   Mean   :0.4496   Mean   :33.31
## 3rd Qu.:105.00   3rd Qu.:36.00   3rd Qu.:0.5800   3rd Qu.:41.00
## Max.   :543.00   Max.   :50.00   Max.   :2.2880   Max.   :67.00
##      class      scored.class      scored.probability
## Min.   :0.0000   Min.   :0.0000   Min.   :0.02323
## 1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.:0.11702
## Median :0.0000   Median :0.0000   Median :0.23999
## Mean   :0.3149   Mean   :0.1768   Mean   :0.30373
## 3rd Qu.:1.0000   3rd Qu.:0.0000   3rd Qu.:0.43093
## Max.   :1.0000   Max.   :1.0000   Max.   :0.94633
```

```
boxplot(output,xlab="variables")
```



Check missing data

```
vis_miss(output)
```



missing data in the file.

Use the table() function to get the raw confusion matrix for this scored dataset (method 1)

```
cf <- table(output[,9:10])
cf
```

```
##      scored.class
## class  0    1
##    0 119    5
##    1  30   27
```

Looking at the matrix above, rows represent actual class values of 0 or 1. Columns represent predicted class values of 0 or 1. So in the top left corner 119 is the number of observations where the class was correctly predicted to be 0. The top right corner shows 5 observations where the class of 0 was incorrectly predicted as 1. Similarly, we have 30 observations of class 1 incorrectly predicted as class 0 and 27 observations of class 1 correctly predicted.

Assuming that 0 is a negative class and 1 is a positive class we have:

TN = 119 FP = 5 FN = 30 TP = 27

Use the table() function to get the raw confusion matrix for this scored dataset (method 2)

```
data <- read.csv("https://raw.githubusercontent.com/johnpannyc/group-1-data-621-assignment-2/master/classification-output-data.csv")
cmatrix <- table(data$class, data$scored.class)
cmatrix
```

```
##
##      0    1
##  0 119    5
##  1   30   27
```

```
Accuracy <- function(df)
{
  names    = c("class", "scored.class")
  cmatrix  = table(df[, names])
  accuracy = (cmatrix[2,2] + cmatrix[1,1]) / (cmatrix[2,2] + cmatrix[1,2] + cmatrix[1,1] + cmatrix[2,1])
  return(round(accuracy, 2))
}
```

```
Accuracy(output)
```

```
## [1] 0.81
```

```
Classification_error_rate <- function(df)
{
  names    = c("class", "scored.class")
  cmatrix  = table(df[, names])
  classification_error_rate = (cmatrix[1,2] + cmatrix[2,1]) / (cmatrix[2,2] + cmatrix[1,2] + cmatrix[1,1] + cmatrix[2,1])
  return(round(classification_error_rate, 2))
}
```

```
Classification_error_rate(output)
```

```
## [1] 0.19
```

```
Precision <- function(df)
{
  names    = c("class", "scored.class")
  cmatrix  = table(df[, names])
  precision = (cmatrix[2,2] / (cmatrix[2,2] + cmatrix[1,2]))
  return(round(precision, 2))
}
```

```
Precision(output)
```

```
## [1] 0.84
```

```
Sensitivity <- function(df)
{
  names      = c("class", "scored.class")
  cmatrix    = table(df[, names])
  sensitivity = cmatrix[2,2] / (cmatrix[2,2] + cmatrix[2,1])
  return(round(sensitivity, 2))
}
```

```
Sensitivity(output)
```

```
## [1] 0.47
```

```
Specificity <- function(df)
{
  names      = c("class", "scored.class")
  cmatrix    = table(df[, names])
  specificity = cmatrix[1,1] / (cmatrix[1,1] + cmatrix[1,2])
  return(round(specificity, 2))
}
```

```
Specificity(output)
```

```
## [1] 0.96
```

```
F1_Score <- function(df)
{
  names      = c("class", "scored.class")
  cmatrix    = table(df[, names])
  precision   = Precision(df)
  sensitivity = Sensitivity(df)
  f1_score   = (2 * precision * sensitivity) / (precision + sensitivity)

  return(round(f1_score, 2))
}
```

```
F1_Score(output)
```

```
## [1] 0.6
```

10. Manually create ROC curve


```

manual_roc <- function(labels, scores){
  labels <- labels[order(scores, decreasing=TRUE)]
  TPR=cumsum(labels)/sum(labels)
  FPR=cumsum(!labels)/sum(!labels)
  df<- data.frame(TPR,FPR)
  dFPR <- c(diff(FPR), 0)
  dTPR <- c(diff(TPR), 0)
  auc <-sum(TPR * dFPR) + sum(dTPR * dFPR)/2
  return(c(df, auc))
}

rc_curve <- manual_roc(output$class,output$scored.probability)

plot(rc_curve[[1]])

```



```
auc <- rc_curve[[2]]
```

Q11 - Using the functions to generate the classification metrics

All metrics were provided as they were calculated. As we will see below using built-in functions makes life easier.

```
Accuracy(output)
```

```
## [1] 0.81
```

```
Classification_error_rate(output)
```

```
## [1] 0.19
```

```
Precision(output)
```

```
## [1] 0.84
```

```
Sensitivity(output)
```

```
## [1] 0.47
```

```
Specificity(output)
```

```
## [1] 0.96
```

```
F1_Score(output)
```

```
## [1] 0.6
```

Q12 - Investigating the caret package

```
if (!"caret" %in% installed.packages()) install.packages(caret)
require(caret)

ls(pos = "package:caret")
```

```

## [1] "anovaScores"          "avNNet"
## [3] "bag"                   "bagControl"
## [5] "bagEarth"             "bagEarthStats"
## [7] "bagFDA"               "best"
## [9] "BoxCoxTrans"          "calibration"
## [11] "caretFuncs"           "caretGA"
## [13] "caretSA"              "caretSBF"
## [15] "caretTheme"           "cforestStats"
## [17] "checkConditionalX"    "checkInstall"
## [19] "checkResamples"       "class2ind"
## [21] "classDist"            "cluster"
## [23] "compare_models"       "confusionMatrix"
## [25] "confusionMatrix.train" "contr.dummy"
## [27] "contr.ltftr"          "createDataPartition"
## [29] "createFolds"          "createModel"
## [31] "createMultiFolds"     "createResample"
## [33] "createTimeSlices"     "ctreeBag"
## [35] "defaultSummary"       "dotPlot"
## [37] "downSample"           "dummyVars"
## [39] "expandParameters"     "expoTrans"
## [41] "extractPrediction"     "extractProb"
## [43] "F_meas"               "featurePlot"
## [45] "filterVarImp"          "findCorrelation"
## [47] "findLinearCombos"     "flatTable"
## [49] "gafs"                 "gafs.default"
## [51] "gafs_initial"         "gafs_lrSelection"
## [53] "gafs_raMutation"      "gafs_rwSelection"
## [55] "gafs_spCrossover"     "gafs_tourSelection"
## [57] "gafs_uCrossover"      "gafsControl"
## [59] "gamFormula"           "gamFuncs"
## [61] "gamScores"            "getModelInfo"
## [63] "getSamplingInfo"      "getTrainPerf"
## [65] "groupKFold"           "hasTerms"
## [67] "icr"                  "index2vec"
## [69] "ipredStats"           "knn3"
## [71] "knn3Train"            "knnreg"
## [73] "knnregTrain"          "ldaBag"
## [75] "ldaFuncs"             "ldaSBF"
## [77] "learing_curve_dat"    "lift"
## [79] "lmFuncs"              "lmSBF"
## [81] "LPH07_1"              "LPH07_2"
## [83] "lrFuncs"              "MAE"
## [85] "maxDissim"            "MeanSD"
## [87] "minDiss"              "mnLogLoss"
## [89] "modelCor"             "modelLookup"
## [91] "multiClassSummary"    "nbBag"
## [93] "nbFuncs"              "nbSBF"
## [95] "nearZeroVar"          "negPredValue"
## [97] "nnetBag"              "nullModel"
## [99] "nzv"                  "oneSE"
## [101] "outcome_conversion"   "panel.calibration"
## [103] "panel.lift"           "panel.lift2"
## [105] "panel.needle"         "pcaNNet"

```



```
## [107] "pickSizeBest"      "pickSizeTolerance"
## [109] "pickVars"          "plot.gafs"
## [111] "plot.rfe"           "plot.train"
## [113] "plotClassProbs"    "plotObsVsPred"
## [115] "plsBag"             "plsda"
## [117] "posPredValue"      "postResample"
## [119] "precision"         "predict.bagEarth"
## [121] "predict.gafs"      "predict.train"
## [123] "predictionFunction" "predictors"
## [125] "preProcess"        "print.train"
## [127] "probFunction"      "progress"
## [129] "prSummary"         "R2"
## [131] "recall"            "resampleHist"
## [133] "resamples"         "resampleSummary"
## [135] "resampleWrapper"   "rfe"
## [137] "rfeControl"        "rfeIter"
## [139] "rfFuncs"           "rfGA"
## [141] "rfSA"              "rfSBF"
## [143] "rfStats"           "RMSE"
## [145] "safs"              "safs_initial"
## [147] "safs_perturb"      "safs_prob"
## [149] "safsControl"       "sbf"
## [151] "sbfControl"        "sbfIter"
## [153] "sensitivity"       "SLC14_1"
## [155] "SLC14_2"           "sortImp"
## [157] "spatialSign"       "specificity"
## [159] "spllda"            "sumDiss"
## [161] "summary.bagEarth"  "svmBag"
## [163] "thresholder"       "tolerance"
## [165] "train"             "trainControl"
## [167] "treebagFuncs"      "treebagGA"
## [169] "treebagSA"         "treebagSBF"
## [171] "twoClassSim"       "twoClassSummary"
## [173] "upSample"          "var_seq"
## [175] "varImp"            "well_numbered"
```

```
?sensitivity
```

```
## starting httpd help server ... done
```

```
?confusionMatrix
?precision
```

Transposing the table so that the actual referenced value (i.e., truth, “class”) is in columns, and the predicted measurement system (i.e. “scored.class”) is in rows

```
df <- data[c("class","scored.class")]
cmatrix.t <- t(table(df))
cmatrix.t
```

```
##           class
## scored.class  0   1
##           0 119  30
##           1   5  27
```

```
str(cmatrix.t)
```

```
## 'table' int [1:2, 1:2] 119 5 30 27
## - attr(*, "dimnames")=List of 2
## ..$ scored.class: chr [1:2] "0" "1"
## ..$ class       : chr [1:2] "0" "1"
```

Comparing the home-made functions and the caret package ones

```
sens.caret <- round(sensitivity(cmatrix.t, positive = rownames(cmatrix)[2]),2)
sens.caret
```

```
## [1] 0.47
```

```
identical(Sensitivity(data), sens.caret)
```

```
## [1] TRUE
```

```
spec.caret <- round(specificity(cmatrix.t, negative = rownames(cmatrix)[1]),2)
spec.caret
```

```
## [1] 0.96
```

```
identical(Specificity(data), spec.caret)
```

```
## [1] TRUE
```

```
cMat.caret <- confusionMatrix(cmatrix.t, positive = "1")
cMat.caret
```

```
## Confusion Matrix and Statistics
##
##           class
## scored.class  0   1
##           0 119  30
##           1   5  27
##
##           Accuracy : 0.8066
##           95% CI : (0.7415, 0.8615)
##           No Information Rate : 0.6851
##           P-Value [Acc > NIR] : 0.0001712
##
##           Kappa : 0.4916
##           McNemar's Test P-Value : 4.976e-05
##
##           Sensitivity : 0.4737
##           Specificity : 0.9597
##           Pos Pred Value : 0.8438
##           Neg Pred Value : 0.7987
##           Prevalence : 0.3149
##           Detection Rate : 0.1492
##           Detection Prevalence : 0.1768
##           Balanced Accuracy : 0.7167
##
##           'Positive' Class : 1
##
```

```
str(cMat.caret)
```

```
## List of 6
## $ positive: chr "1"
## $ table : 'table' int [1:2, 1:2] 119 5 30 27
## .. attr(*, "dimnames")=List of 2
## .. ..$ scored.class: chr [1:2] "0" "1"
## .. ..$ class : chr [1:2] "0" "1"
## $ overall : Named num [1:7] 0.807 0.492 0.741 0.861 0.685 ...
## .. attr(*, "names")= chr [1:7] "Accuracy" "Kappa" "AccuracyLower" "AccuracyUpper" ...
## $ byClass : Named num [1:11] 0.474 0.96 0.844 0.799 0.844 ...
## .. attr(*, "names")= chr [1:11] "Sensitivity" "Specificity" "Pos Pred Value" "Neg Pred Value" ...
## $ mode : chr "sens_spec"
## $ dots : list()
## - attr(*, "class")= chr "confusionMatrix"
```

```
prec.caret <- round(precision(cmatrix.t, relevant = "1"),2)
prec.caret
```

```
## [1] 0.84
```

```
identical(Precision(data), prec.caret)
```

```
## [1] TRUE
```

```
acc.caret <- round(cMat.caret$overall[1],2)  
acc.caret
```


```
## Accuracy  
##      0.81
```

```
identical(Accuracy(data), acc.caret)  ## same value, but fail to match with identical function
```

```
## [1] FALSE
```

13. pROC Package Let us try the pROC package.

```
roc(output$class, output$scored.probability, levels=c(0,1), percent=TRUE, plot=TRUE, ci=TRUE)
```



```
##  
## Call:  
## roc.default(response = output$class, predictor = output$scored.probability,      levels = c(0,  
1), percent = TRUE, ci = TRUE, plot = TRUE)  
##  
## Data: output$scored.probability in 124 controls (output$class 0) < 57 cases (output$class 1).  
## Area under the curve: 85.03%  
## 95% CI: 79.05%-91.01% (DeLong)
```